

November 18, 2024

## 1 Data Collection and Preparation

```
[ ]: # Import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

import seaborn as sns

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, \
    f1_score, roc_auc_score, classification_report, confusion_matrix
```

```
[ ]: data = pd.read_excel('account_churn_project.xlsx')
```

```
[ ]: print(data.shape)
```

```
(10127, 21)
```

```
[ ]: print(data.head())
```

	identification	churn_flag	age	gender	number_dependants	\
0	768805383	Existing Customer	45	M	3	
1	818770008	Existing Customer	49	F	5	
2	713982108	Existing Customer	51	M	3	
3	769911858	Existing Customer	40	F	4	
4	709106358	Existing Customer	40	M	3	

	education	civil_status	income	account_category	account_age	\
0	High School	Married	\$60K - \$80K	Blue	39	
1	Graduate	Single	Less than \$40K	Blue	44	
2	Graduate	Married	\$80K - \$120K	Blue	36	
3	High School	Unknown	Less than \$40K	Blue	34	

4	Uneducated	Married	\$60K - \$80K	Blue	21
---	------------	---------	---------------	------	----

	...	inactivity	past_contacts	card_Limit	balance	open_to_use	\
0	...	1	3	12691	777	11914	
1	...	1	2	8256	864	7392	
2	...	1	0	3418	0	3418	
3	...	4	1	3313	2517	796	
4	...	1	0	4716	0	4716	

	change_per_quarter_amount	total_transaction_amount	\
0	1335	1144	
1	1541	1291	
2	2594	1887	
3	1405	1171	
4	2175	816	

	total_transaction_count	change_per_quarter_quantity	average_use
0	42	1625	61
1	33	3714	105
2	20	2333	0
3	20	2333	76
4	28	25	0

[5 rows x 21 columns]

```
[ ]: # 1. Check data types
print("Data types:")
print(data.dtypes)
```

```
Data types:
identification      int64
churn_flag           object
age                 int64
gender              object
number_dependants    int64
education            object
civil_status         object
income              object
account_category     object
account_age          int64
total_num_services   int64
inactivity           int64
past_contacts        int64
card_Limit           int64
balance             int64
open_to_use          int64
change_per_quarter_amount int64
total_transaction_amount int64
```

```
total_transaction_count      int64
change_per_quarter_quantity  int64
average_use                  int64
dtype: object
```

```
[ ]: # 2. Check for missing values
print("\nMissing values by column:")
print(data.isnull().sum())
```

```
Missing values by column:
identification      0
churn_flag          0
age                 0
gender              0
number_dependants   0
education           0
civil_status        0
income              0
account_category    0
account_age         0
total_num_services  0
inactivity          0
past_contacts       0
card_Limit          0
balance             0
open_to_use         0
change_per_quarter_amount  0
total_transaction_amount  0
total_transaction_count  0
change_per_quarter_quantity  0
average_use         0
dtype: int64
```

```
[ ]: # 3. Descriptive statistics
print("\nDescriptive statistics:")
print(data.describe())
```

```
Descriptive statistics:
```

	identification	age	number_dependants	account_age \
count	1.012700e+04	10127.000000	10127.000000	10127.000000
mean	7.391776e+08	46.325960	2.346203	35.928409
std	3.690378e+07	8.016814	1.298908	7.986416
min	7.080821e+08	26.000000	0.000000	13.000000
25%	7.130368e+08	41.000000	1.000000	31.000000
50%	7.179264e+08	46.000000	2.000000	36.000000
75%	7.731435e+08	52.000000	3.000000	40.000000

max	8.283431e+08	73.000000	5.000000	56.000000
-----	--------------	-----------	----------	-----------

	total_num_services	inactivity	past_contacts	card_Limit \
count	10127.000000	10127.000000	10127.000000	10127.000000
mean	3.812580	2.341167	2.455317	9280.019552
std	1.554408	1.010622	1.106225	9013.924409
min	1.000000	0.000000	0.000000	1439.000000
25%	3.000000	2.000000	2.000000	2787.000000
50%	4.000000	2.000000	2.000000	5363.000000
75%	5.000000	3.000000	3.000000	13576.000000
max	6.000000	6.000000	6.000000	34516.000000

	balance	open_to_use	change_per_quarter_amount \
count	10127.000000	10127.000000	10127.000000
mean	1162.814061	7988.004345	689.673645
std	814.987335	9054.851654	297.521767
min	0.000000	3.000000	0.000000
25%	359.000000	1463.500000	579.000000
50%	1276.000000	4199.000000	715.000000
75%	1784.000000	11456.000000	844.000000
max	2517.000000	34516.000000	3397.000000

	total_transaction_amount	total_transaction_count \
count	10127.000000	10127.000000
mean	4404.086304	64.858695
std	3397.129254	23.472570
min	510.000000	10.000000
25%	2155.500000	45.000000
50%	3899.000000	67.000000
75%	4741.000000	81.000000
max	18484.000000	139.000000

	change_per_quarter_quantity	average_use
count	10127.000000	10127.000000
mean	591.557322	249.842500
std	323.649708	272.424923
min	0.000000	0.000000
25%	438.500000	3.500000
50%	655.000000	132.000000
75%	786.000000	463.000000
max	3714.000000	999.000000

```
[ ]: # 4. Check for duplicates
print("\nNumber of duplicates:")
print(data.duplicated().sum())
```

Number of duplicates:

0

```
[ ]: # 5. Check value ranges
# Checking age
print("\nChecking inconsistent ages (less than 0 or greater than 120):")
print(data[(data['age'] < 0) | (data['age'] > 120)])

# Ensure all of this columns are positive or zero)
# List of columns you want to verify for int64 type and positive/zero values
columns = ['number_dependants', 'account_age', 'card_Limit', 'balance',
           'open_to_use',
           'total_transaction_amount', 'total_transaction_count',
           'change_per_quarter_amount',
           'change_per_quarter_quantity', 'average_use']

for column in columns:
    print(f"\nChecking {column} for negative values:")
    print(data[data[column] < 0])
```

Checking inconsistent ages (less than 0 or greater than 120):

Empty DataFrame

Columns: [identification, churn\_flag, age, gender, number\_dependants, education, civil\_status, income, account\_category, account\_age, total\_num\_services, inactivity, past\_contacts, card\_Limit, balance, open\_to\_use, change\_per\_quarter\_amount, total\_transaction\_amount, total\_transaction\_count, change\_per\_quarter\_quantity, average\_use]

Index: []

[0 rows x 21 columns]

Checking number\_dependants for negative values:

Empty DataFrame

Columns: [identification, churn\_flag, age, gender, number\_dependants, education, civil\_status, income, account\_category, account\_age, total\_num\_services, inactivity, past\_contacts, card\_Limit, balance, open\_to\_use, change\_per\_quarter\_amount, total\_transaction\_amount, total\_transaction\_count, change\_per\_quarter\_quantity, average\_use]

Index: []

[0 rows x 21 columns]

Checking account\_age for negative values:

Empty DataFrame

Columns: [identification, churn\_flag, age, gender, number\_dependants, education, civil\_status, income, account\_category, account\_age, total\_num\_services, inactivity, past\_contacts, card\_Limit, balance, open\_to\_use, change\_per\_quarter\_amount, total\_transaction\_amount, total\_transaction\_count,

```
change_per_quarter_quantity, average_use]
Index: []
```

```
[0 rows x 21 columns]
```

Checking card\_Limit for negative values:

Empty DataFrame

Columns: [identification, churn\_flag, age, gender, number\_dependants, education, civil\_status, income, account\_category, account\_age, total\_num\_services, inactivity, past\_contacts, card\_Limit, balance, open\_to\_use, change\_per\_quarter\_amount, total\_transaction\_amount, total\_transaction\_count, change\_per\_quarter\_quantity, average\_use]

Index: []

```
[0 rows x 21 columns]
```

Checking balance for negative values:

Empty DataFrame

Columns: [identification, churn\_flag, age, gender, number\_dependants, education, civil\_status, income, account\_category, account\_age, total\_num\_services, inactivity, past\_contacts, card\_Limit, balance, open\_to\_use, change\_per\_quarter\_amount, total\_transaction\_amount, total\_transaction\_count, change\_per\_quarter\_quantity, average\_use]

Index: []

```
[0 rows x 21 columns]
```

Checking open\_to\_use for negative values:

Empty DataFrame

Columns: [identification, churn\_flag, age, gender, number\_dependants, education, civil\_status, income, account\_category, account\_age, total\_num\_services, inactivity, past\_contacts, card\_Limit, balance, open\_to\_use, change\_per\_quarter\_amount, total\_transaction\_amount, total\_transaction\_count, change\_per\_quarter\_quantity, average\_use]

Index: []

```
[0 rows x 21 columns]
```

Checking total\_transaction\_amount for negative values:

Empty DataFrame

Columns: [identification, churn\_flag, age, gender, number\_dependants, education, civil\_status, income, account\_category, account\_age, total\_num\_services, inactivity, past\_contacts, card\_Limit, balance, open\_to\_use, change\_per\_quarter\_amount, total\_transaction\_amount, total\_transaction\_count, change\_per\_quarter\_quantity, average\_use]

Index: []

```
[0 rows x 21 columns]
```

Checking total\_transaction\_count for negative values:

Empty DataFrame

Columns: [identification, churn\_flag, age, gender, number\_dependants, education, civil\_status, income, account\_category, account\_age, total\_num\_services, inactivity, past\_contacts, card\_limit, balance, open\_to\_use, change\_per\_quarter\_amount, total\_transaction\_amount, total\_transaction\_count, change\_per\_quarter\_quantity, average\_use]

Index: []

[0 rows x 21 columns]

Checking change\_per\_quarter\_amount for negative values:

Empty DataFrame

Columns: [identification, churn\_flag, age, gender, number\_dependants, education, civil\_status, income, account\_category, account\_age, total\_num\_services, inactivity, past\_contacts, card\_limit, balance, open\_to\_use, change\_per\_quarter\_amount, total\_transaction\_amount, total\_transaction\_count, change\_per\_quarter\_quantity, average\_use]

Index: []

[0 rows x 21 columns]

Checking change\_per\_quarter\_quantity for negative values:

Empty DataFrame

Columns: [identification, churn\_flag, age, gender, number\_dependants, education, civil\_status, income, account\_category, account\_age, total\_num\_services, inactivity, past\_contacts, card\_limit, balance, open\_to\_use, change\_per\_quarter\_amount, total\_transaction\_amount, total\_transaction\_count, change\_per\_quarter\_quantity, average\_use]

Index: []

[0 rows x 21 columns]

Checking average\_use for negative values:

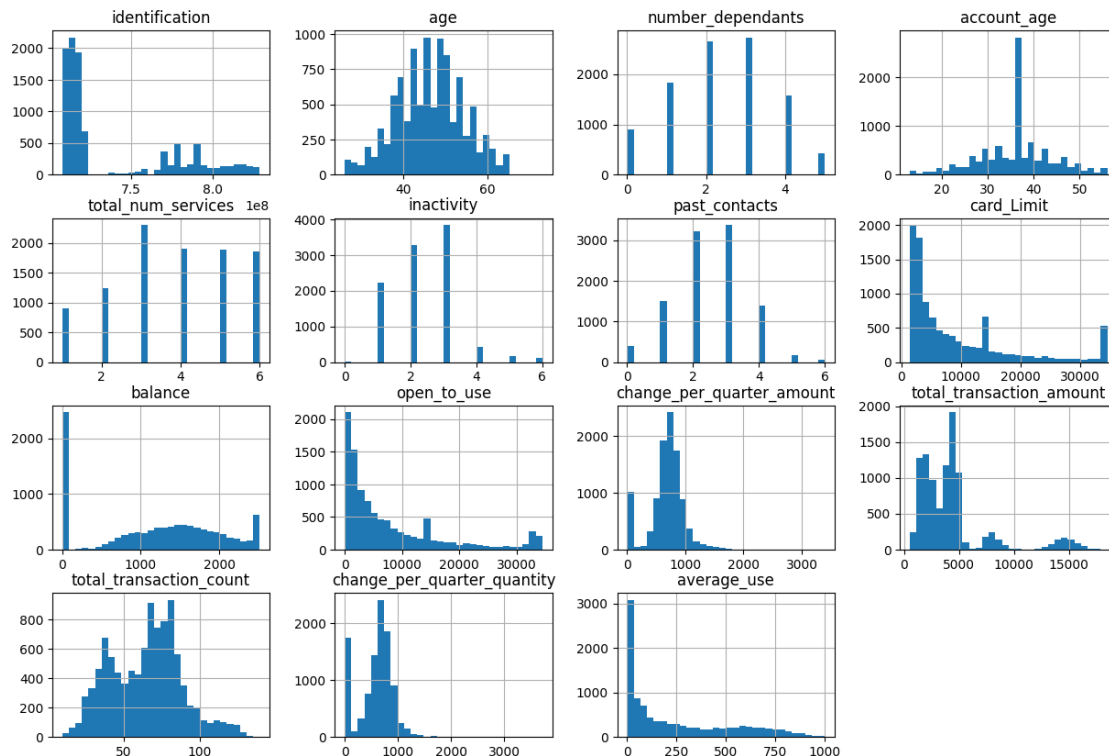
Empty DataFrame

Columns: [identification, churn\_flag, age, gender, number\_dependants, education, civil\_status, income, account\_category, account\_age, total\_num\_services, inactivity, past\_contacts, card\_limit, balance, open\_to\_use, change\_per\_quarter\_amount, total\_transaction\_amount, total\_transaction\_count, change\_per\_quarter\_quantity, average\_use]

Index: []

[0 rows x 21 columns]

```
[ ]: data.hist(bins=30, figsize=(15, 10))  
      plt.show()
```



```
[ ]: # Map churn_flag values to binary (1 for Existing Customer, 0 for Attrited
      ↪Customer)
data['churn_flag'] = data['churn_flag'].map({'Existing Customer': 0, 'Attrited_
      ↪Customer': 1})

# Map gender values ('M' becomes 0, 'F' becomes 1)
data['gender'] = data['gender'].map({'M': 0, 'F': 1})

# Transform categorical data into dummy/indicator variables
df_encoded = pd.get_dummies(data)

# Display the first 5 rows of the encoded DataFrame
print(df_encoded.head())
```

	identification	churn_flag	age	gender	number_dependants	account_age	\
0	768805383	0	45	0	3	39	
1	818770008	0	49	1	5	44	
2	713982108	0	51	0	3	36	
3	769911858	0	40	1	4	34	
4	709106358	0	40	0	3	21	

	total_num_services	inactivity	past_contacts	card_Limit	...	\
0	5	1	3	12691	...	



1	6	1	2	8256	...
2	4	1	0	3418	...
3	3	4	1	3313	...
4	5	1	0	4716	...

	income_>120K +	income_>40K - >60K	income_>60K - >80K	\
0	False	False	True	
1	False	False	False	
2	False	False	False	
3	False	False	False	
4	False	False	True	

	income_>80K - >120K	income_Less than >40K	income_Unknown	\
0	False	False	False	
1	False	True	False	
2	True	False	False	
3	False	True	False	
4	False	False	False	

	account_category_Blue	account_category_Gold	account_category_Platinum	\
0	True	False	False	
1	True	False	False	
2	True	False	False	
3	True	False	False	
4	True	False	False	

	account_category_Silver
0	False
1	False
2	False
3	False
4	False

[5 rows x 38 columns]

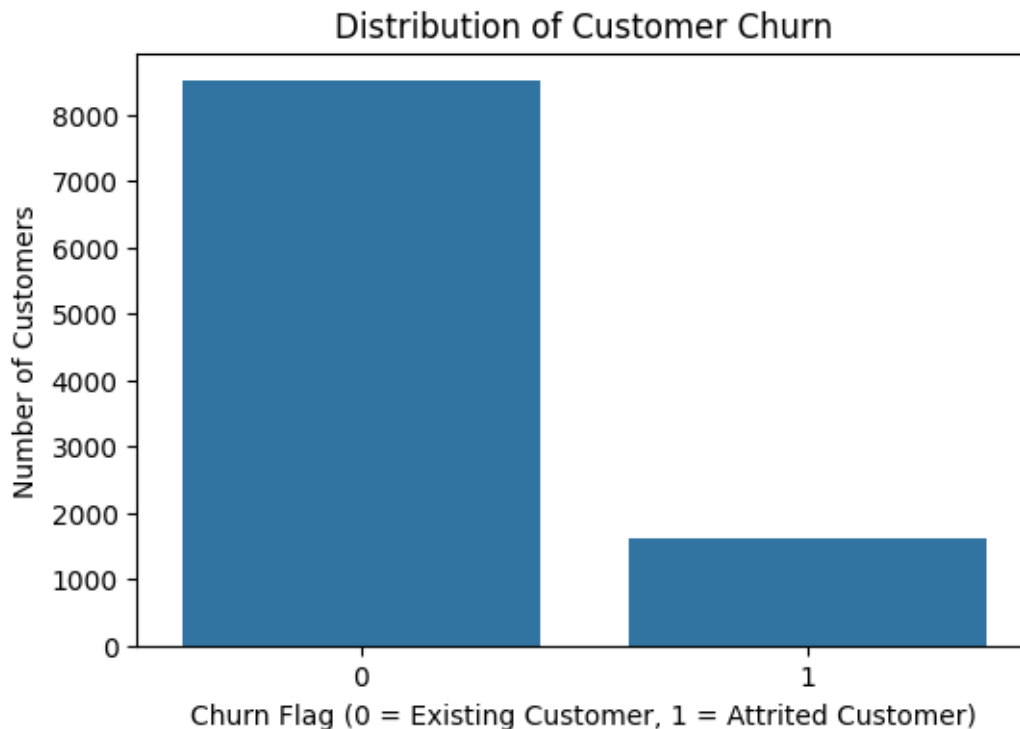
```
[ ]: # Create new features
df_encoded['card_utilization'] = df_encoded['balance'] /
    df_encoded['card_Limit']
```

```
[ ]: df_encoded.to_excel('modified_file.xlsx', index=False)
```

## 2 Exploratory Data Analysis


### 3 Visualizing Customer Churn with a Count Plot

```
[ ]: plt.figure(figsize=(6,4))
      # Create a count plot to display the distribution of churn flag Customer
      ax = sns.countplot(x='churn_flag',data=df_encoded)
      plt.title('Distribution of Customer Churn')
      plt.xlabel('Churn Flag (0 = Existing Customer, 1 = Attrited Customer)')
      plt.ylabel('Number of Customers')
      plt.show()
```



## 4 Visualizing the Correlation Matrix of Features

```
[ ]: plt.figure(figsize=(12,8))
      corr_matrix = df_encoded.corr()
      sns.heatmap(corr_matrix, annot=False, cmap='coolwarm', linewidths=0.5)
      plt.title('Correlation Matrix of Features')
      plt.show()
```




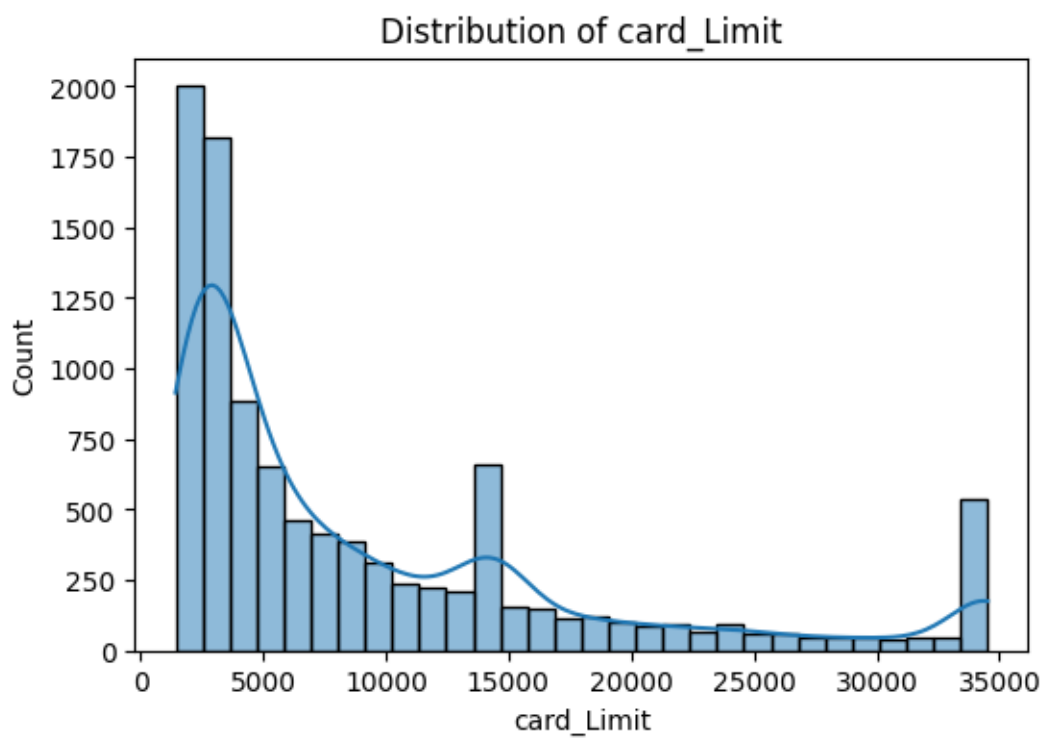
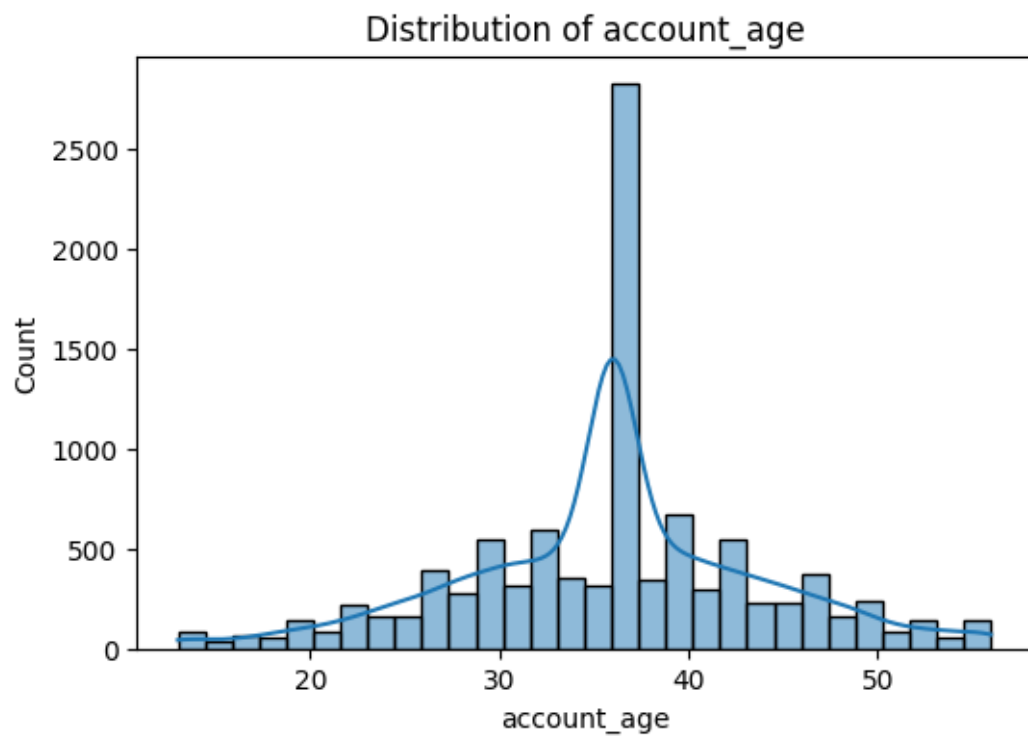
```
[ ]: corr = df_encoded.corr()
target_corr = corr['churn_flag'].abs().sort_values(ascending=False)
print(target_corr.head(10))
```



```
churn_flag          1.000000
total_transaction_count  0.371403
balance             0.263053
past_contacts       0.204491
change_per_quarter_quantity  0.195368
card_utilization    0.177722
total_transaction_amount  0.168598
average_use         0.163300
inactivity          0.152449
total_num_services  0.150005
Name: churn_flag, dtype: float64
```

```
[ ]: # Visualize numerical distributions
num_vars = ['age', 'account_age', 'card_limit', 'balance', 'total_transaction_amount']
for col in num_vars:
```

```
plt.figure(figsize=(6, 4))
sns.histplot(df_encoded[col], kde=True, bins=30)
plt.title(f'Distribution of {col}')
plt.show()
```







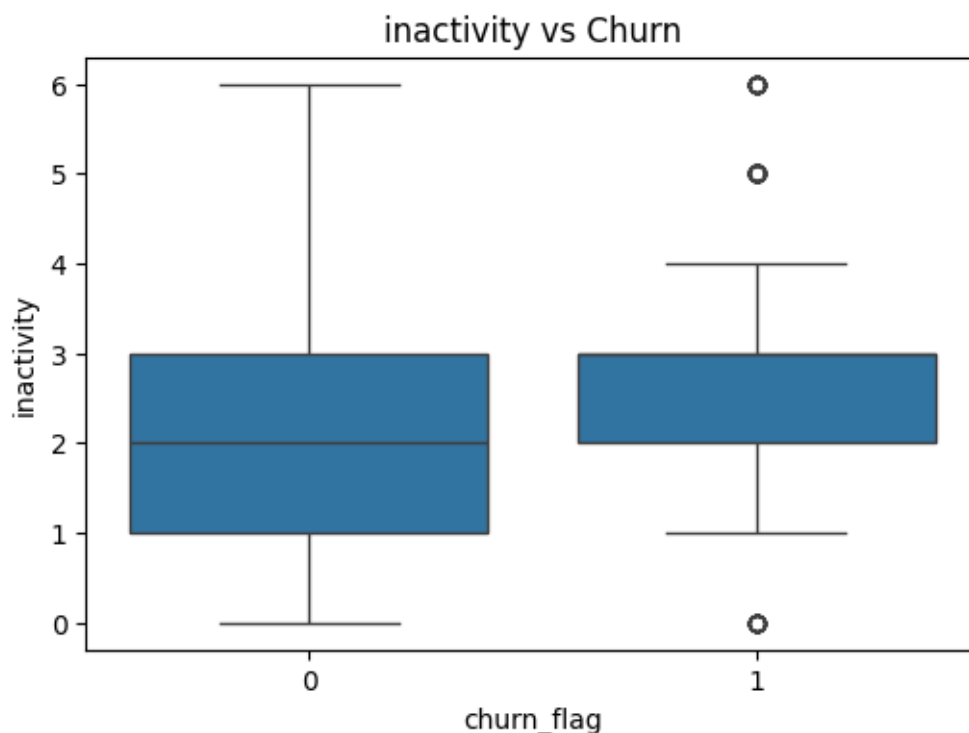
```
[ ]: # Boxplots for relationships


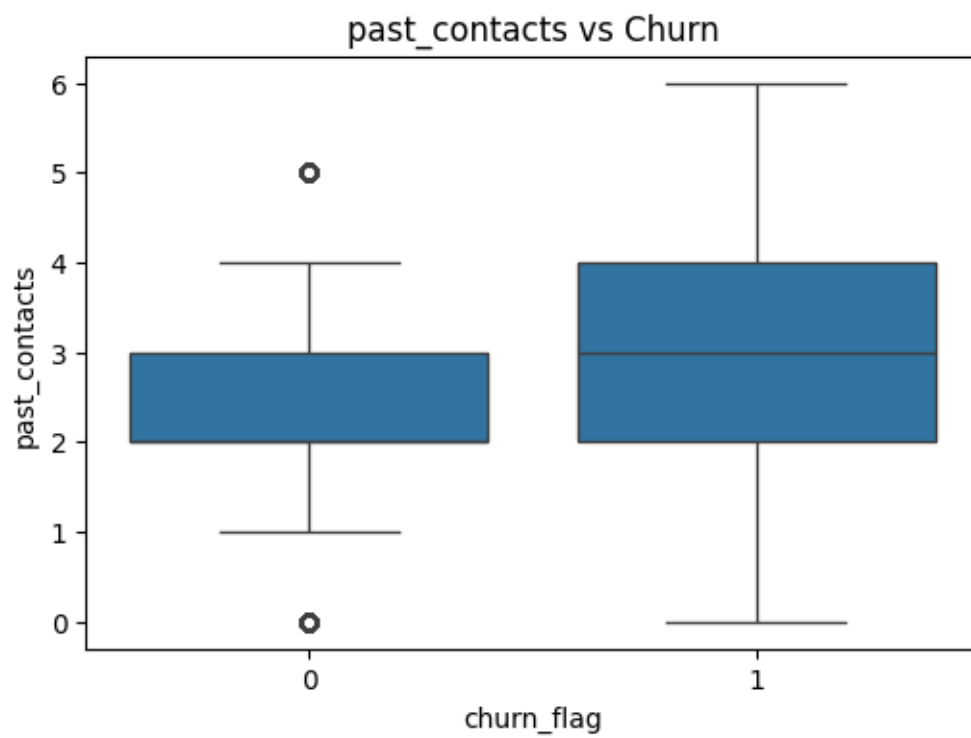
key_features = [
    'inactivity', 'past_contacts', 'balance', 'total_transaction_amount', 'total_transaction_count',
    'age', 'balance', 'card_Limit',
    'total_transaction_amount', 'total_transaction_count', 'past_contacts', 'change_per_quarter_qu
#all_columns = df_encoded.columns
#columns_to_plot = [col for col in all_columns if col not in ['churn_flag',
    'identification']]

#for feature in columns_to_plot:
    # plt.figure(figsize=(6, 4))
    # sns.boxplot(x='churn_flag', y=feature, data=df_encoded)
    # plt.title(f'{feature} vs Churn')
    # plt.show()

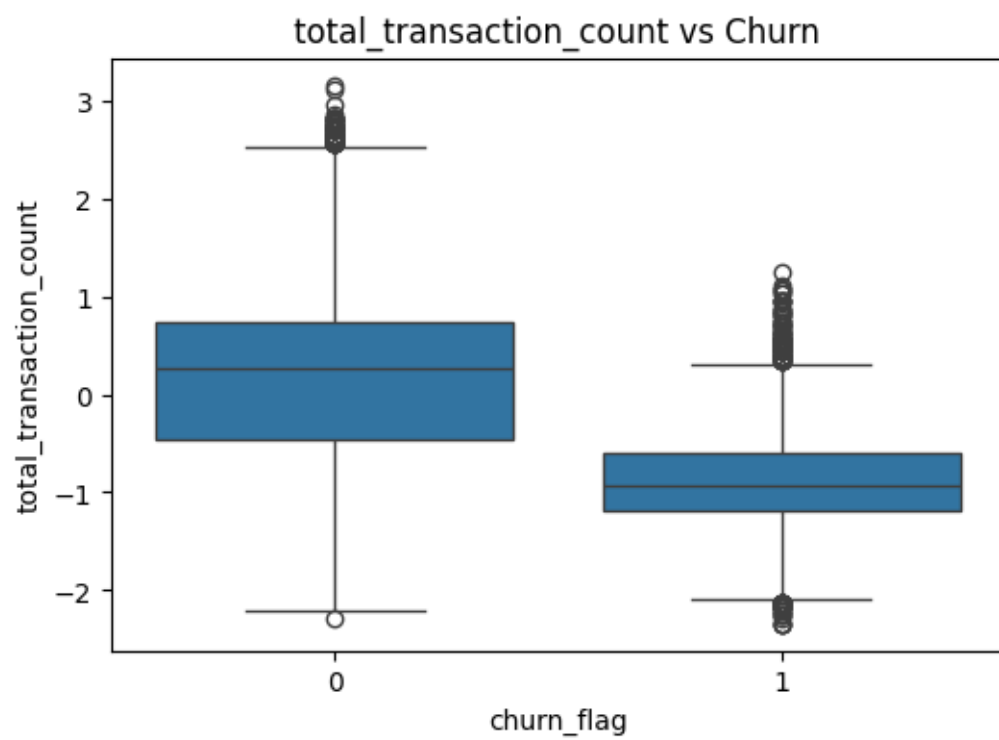

print('0 = existing customers; 1 = attrited customers')
for feature in key_features:
    plt.figure(figsize=(6, 4))
    sns.boxplot(x='churn_flag', y=feature, data=df_encoded)
    plt.title(f'{feature} vs Churn')
    plt.show()
```

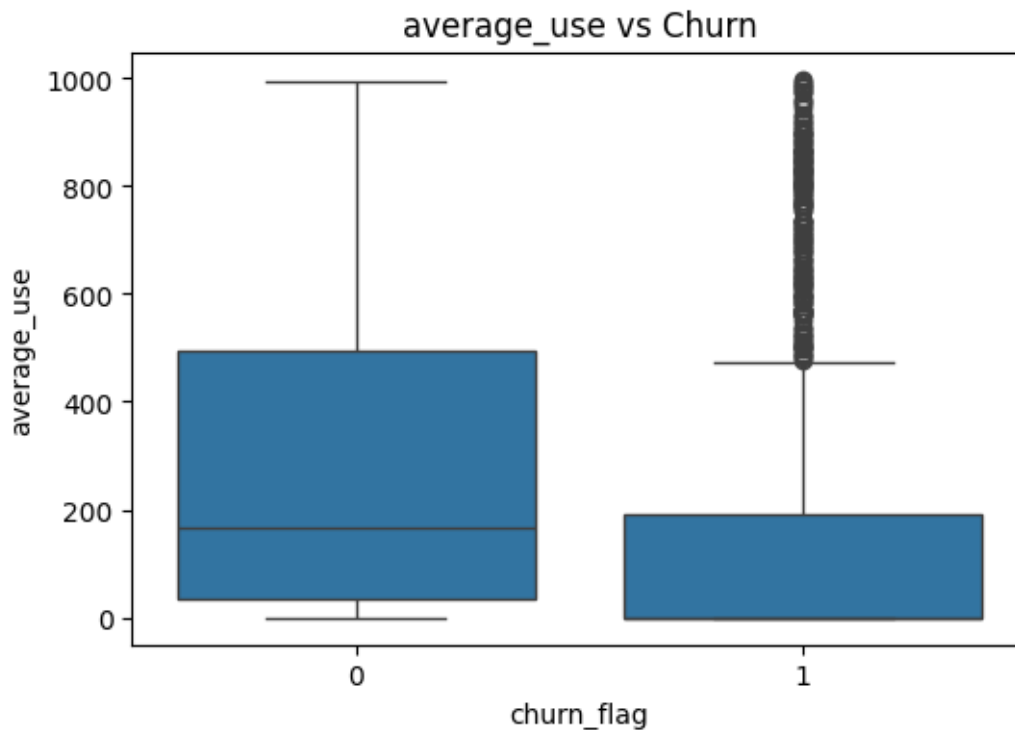
0 = existing customers; 1 = attrited customers











```
[ ]: # Overlay histograms or KDE plots for churned vs. existing customers
#key_features = ['balance',
↳ 'card_Limit', 'total_transaction_count', 'past_contacts', 'change_per_quarter_quantity', 'average_use']
key_features = ['total_num_services',
↳ 'balance', 'total_transaction_amount', 'total_transaction_count', 'average_use']

print('0 = existing customers; 1 = attrited customers')
for feature in key_features:
    #for feature in columns_to_plot:
        plt.figure(figsize=(8, 6))

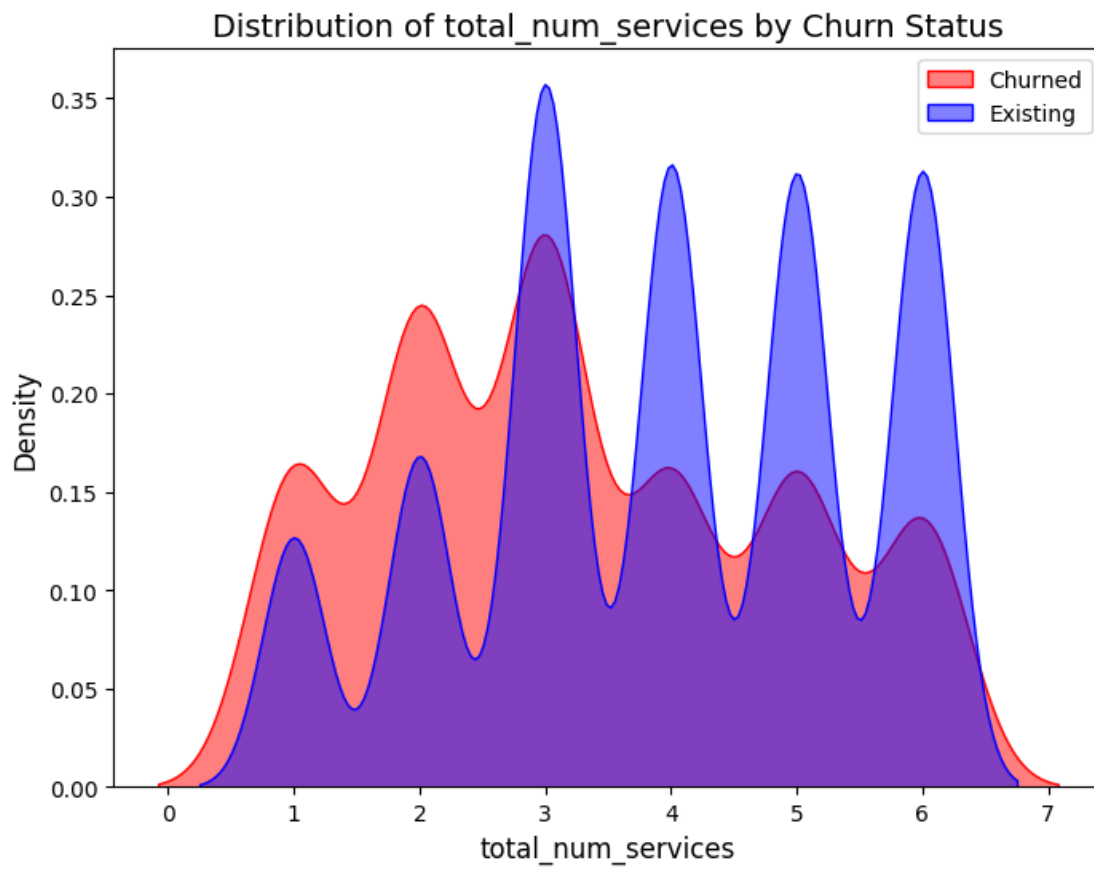
        # KDE plot for churned customers
        sns.kdeplot(data=df_encoded[df_encoded['churn_flag'] == 1][feature],
                    label='Churned', color='red', fill=True, alpha=0.5)


        # KDE plot for existing customers
        sns.kdeplot(data=df_encoded[df_encoded['churn_flag'] == 0][feature],
                    label='Existing', color='blue', fill=True, alpha=0.5)


        plt.title(f'Distribution of {feature} by Churn Status', fontsize=14)
        plt.xlabel(feature, fontsize=12)
```


```
plt.ylabel('Density', fontsize=12)
plt.legend()
plt.show()
```

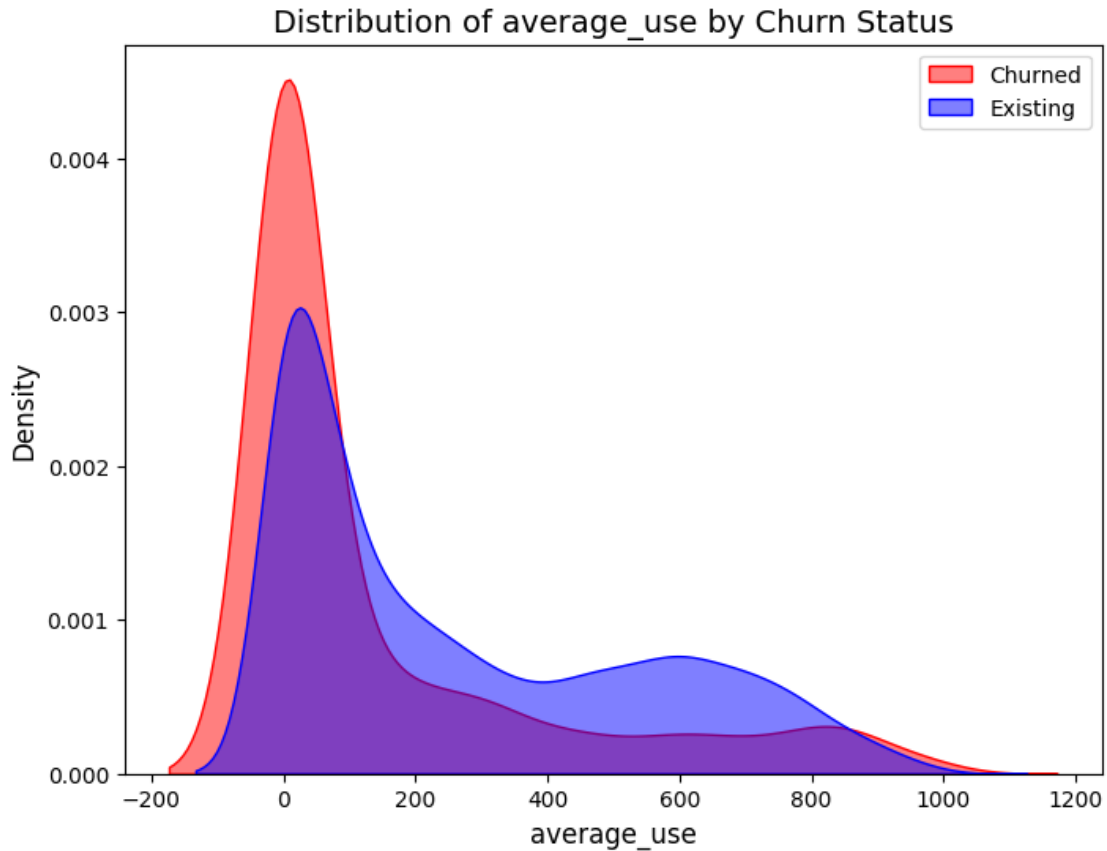
0 = existing customers; 1 = attrited customers











```
[ ]: # Scale numerical features
scaler = StandardScaler()
num_cols = ['age', 'number_dependants', 'account_age', 'card_Limit', 'balance',
            'total_transaction_amount', 'total_transaction_count']
df_encoded[num_cols] = scaler.fit_transform(df_encoded[num_cols])

print(df_encoded.head())
```

	identification	churn_flag	age	gender	number_dependants	\
0	768805383	0	-0.165406	0	0.503368	
1	818770008	0	0.333570	1	2.043199	
2	713982108	0	0.583058	0	0.503368	
3	769911858	0	-0.789126	1	1.273283	
4	709106358	0	-0.789126	0	0.503368	

	account_age	total_num_services	inactivity	past_contacts	card_Limit	\
0	0.384621	5	1	3	0.378431	
1	1.010715	6	1	2	-0.113610	
2	0.008965	4	1	0	-0.650361	
3	-0.241473	3	4	1	-0.662011	

4	-1.869317	5	1	0	-0.506355
---	-----------	---	---	---	-----------

	...	income_\$40K - \$60K	income_\$60K - \$80K	income_\$80K - \$120K	\
0	...	False	True	False	
1	...	False	False	False	
2	...	False	False	True	
3	...	False	False	False	
4	...	False	True	False	

		income_Less than \$40K	income_Unknown	account_category_Blue	\
0		False	False	True	
1		True	False	True	
2		False	False	True	
3		True	False	True	
4		False	False	True	

		account_category_Gold	account_category_Platinum	account_category_Silver	\
0		False	False	False	
1		False	False	False	
2		False	False	False	
3		False	False	False	
4		False	False	False	

	card_utilization
0	0.061224
1	0.104651
2	0.000000
3	0.759734
4	0.000000

[5 rows x 39 columns]

## 5 Feature Selection

```
[ ]: corr = df_encoded.corr()
target_corr = corr['churn_flag'].abs().sort_values(ascending=False)
print(target_corr.head(10))
```

churn_flag	1.000000
total_transaction_count	0.371403
balance	0.263053
past_contacts	0.204491
change_per_quarter_quantity	0.195368
card_utilization	0.177722
total_transaction_amount	0.168598
average_use	0.163300
inactivity	0.152449



```
total_num_services      0.150005
Name: churn_flag, dtype: float64
```

## 6 Model Building

```
[ ]: # Split the dataset into training and testing subsets for model development and
      ↪ validation.
      # Features and target
      X = df_encoded.drop(columns=['churn_flag', 'identification'])
      y = df_encoded['churn_flag']

      # Split the data
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
      ↪ random_state=42, stratify=y)
      from sklearn.preprocessing import StandardScaler

      # Assuming 'num_cols' contains the numerical features
      scaler = StandardScaler()
      X_train[num_cols] = scaler.fit_transform(X_train[num_cols])
      X_test[num_cols] = scaler.transform(X_test[num_cols])
```

## 7 Logistic Regression

```
[ ]: # Logistic Regression
      log_reg = LogisticRegression(class_weight='balanced',
      ↪ random_state=42, max_iter=1000)
      log_reg.fit(X_train, y_train)

      # Predictions and evaluation
      y_pred_log = log_reg.predict(X_test)
      print('Logistic Regression Metrics:')
      print(classification_report(y_test, y_pred_log))

      cm = confusion_matrix(y_test, y_pred_log)
      sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
```

Logistic Regression Metrics:

	precision	recall	f1-score	support
0	0.96	0.84	0.90	2551
1	0.50	0.81	0.62	488
accuracy			0.84	3039
macro avg	0.73	0.83	0.76	3039
weighted avg	0.88	0.84	0.85	3039

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:469:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```


Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(  
[ ]: <Axes: >
```



## 8 K-Nearest Neighbors

```
[ ]: # KNN  
from imblearn.over_sampling import SMOTE  
  
# Instantiate SMOTE  
smote = SMOTE(random_state=42)
```

```

# Resample the training data
X_train_resampled, y_train_resampled = smote.fit_resample(X_train, y_train)

knn = KNeighborsClassifier()
knn_params = {'n_neighbors': [3, 5, 7], 'weights': ['uniform', 'distance']}
knn_grid = GridSearchCV(knn, knn_params, cv=3, scoring='f1')
knn_grid.fit(X_train_resampled, y_train_resampled)

# Best KNN
best_knn = knn_grid.best_estimator_
y_pred_knn = best_knn.predict(X_test)
print('KNN Metrics:')
print(classification_report(y_test, y_pred_knn))

```

KNN Metrics:

	precision	recall	f1-score	support
0	0.89	0.80	0.84	2551
1	0.32	0.50	0.39	488
accuracy			0.75	3039
macro avg	0.61	0.65	0.62	3039
weighted avg	0.80	0.75	0.77	3039

## 9 Random Forest

```

[ ]: # Random Forest
rf = RandomForestClassifier(class_weight='balanced', random_state=42)
rf_params = {'n_estimators': [100, 200], 'max_depth': [10, 20, None]}
rf_grid = GridSearchCV(rf, rf_params, cv=3, scoring='f1')
rf_grid.fit(X_train, y_train)

# Best Random Forest
best_rf = rf_grid.best_estimator_
y_pred_rf = best_rf.predict(X_test)
print('Random Forest Metrics:')
print(classification_report(y_test, y_pred_rf))

```


Random Forest Metrics:

	precision	recall	f1-score	support
0	0.97	0.96	0.97	2551
1	0.81	0.83	0.82	488
accuracy			0.94	3039
macro avg	0.89	0.90	0.89	3039

weighted avg      0.94      0.94      0.94      3039

```
[ ]: # Feature importance from Random Forest
importances = best_rf.feature_importances_
feature_names = X.columns
feature_importance_df = pd.DataFrame({'Feature': feature_names, 'Importance':
    ↳ importances})
feature_importance_df = feature_importance_df.sort_values(by='Importance',
    ↳ ascending=False)

# Plot
plt.figure(figsize=(10, 6))
sns.barplot(x='Importance', y='Feature', data=feature_importance_df)
plt.title('Feature Importance from Random Forest')
plt.show()
```



```
[ ]: rf = RandomForestClassifier(random_state=42)
rf.fit(X_train, y_train)
feature_importances = pd.DataFrame({'Feature': X.columns,
    ↳ 'Importance': rf.feature_importances_}).
    ↳ sort_values(by='Importance', ascending=False)
print(feature_importances)
```

	Feature	Importance
11	total_transaction_amount	0.190330
12	total_transaction_count	0.163504
8	balance	0.093679

4	total_num_services	0.068026
13	change_per_quarter_quantity	0.061804
36	card_utilization	0.058197
10	change_per_quarter_amount	0.054437
14	average_use	0.045589
7	card_Limit	0.035678
9	open_to_use	0.034838
0	age	0.031312
6	past_contacts	0.029549
3	account_age	0.025280
5	inactivity	0.025108
2	number_dependants	0.014094
1	gender	0.009217
23	civil_status_Married	0.005549
24	civil_status_Single	0.004867
17	education_Graduate	0.004137
30	income_Less than \$40K	0.003749
21	education_Unknown	0.003446
28	income_\$60K - \$80K	0.003399
18	education_High School	0.003354
27	income_\$40K - \$60K	0.003321
20	education_Uneducated	0.003131
25	civil_status_Unknown	0.003085
29	income_\$80K - \$120K	0.003007
15	education_College	0.002658
31	income_Unknown	0.002453
22	civil_status_Divorced	0.002124
26	income_\$120K +	0.001976
19	education_Post-Graduate	0.001905
35	account_category_Silver	0.001879
32	account_category_Blue	0.001876
16	education_Doctorate	0.001836
33	account_category_Gold	0.001376
34	account_category_Platinum	0.000231

## 10 Model Evaluation

```
[ ]: from sklearn.metrics import roc_auc_score, roc_curve, auc

# Logistic Regression
y_pred_log_proba = log_reg.predict_proba(X_test)[:, 1] # Probabilities for the
               ↪ positive class
log_reg_auc = roc_auc_score(y_test, y_pred_log_proba)
print(f"Logistic Regression AUC: {log_reg_auc:.4f}")

# K-Nearest Neighbors
```

```

y_pred_knn_proba = best_knn.predict_proba(X_test)[: , 1] # Probabilities for the
↳the positive class
knn_auc = roc_auc_score(y_test, y_pred_knn_proba)
print(f"KNN AUC: {knn_auc:.4f}")

# Random Forest
y_pred_rf_proba = best_rf.predict_proba(X_test)[: , 1] # Probabilities for the
↳positive class
rf_auc = roc_auc_score(y_test, y_pred_rf_proba)
print(f"Random Forest AUC: {rf_auc:.4f}")

# Plot ROC Curve for all models
plt.figure(figsize=(10, 8))

# Logistic Regression ROC
fpr_log, tpr_log, _ = roc_curve(y_test, y_pred_log_proba)
plt.plot(fpr_log, tpr_log, label=f"Logistic Regression (AUC = {log_reg_auc:.
↳4f})")

# KNN ROC
fpr_knn, tpr_knn, _ = roc_curve(y_test, y_pred_knn_proba)
plt.plot(fpr_knn, tpr_knn, label=f"KNN (AUC = {knn_auc:.4f})")

# Random Forest ROC
fpr_rf, tpr_rf, _ = roc_curve(y_test, y_pred_rf_proba)
plt.plot(fpr_rf, tpr_rf, label=f"Random Forest (AUC = {rf_auc:.4f})")

# Plot settings
plt.plot([0, 1], [0, 1], 'k--', label="Random Chance (AUC = 0.5000)")
plt.title("ROC Curve for Models")
plt.xlabel("False Positive Rate (FPR)")
plt.ylabel("True Positive Rate (TPR)")
plt.legend(loc="lower right")
plt.grid()
plt.show()

```

Logistic Regression AUC: 0.9069  
 KNN AUC: 0.6896  
 Random Forest AUC: 0.9752

